

Exercises 1

Coauthored by: Teng, Chao & Jonas Karl, Böcker

Problem 1

The nonparallel version of computing pi by Monte Carlo.

```
void compute_pi_monte_carlo(int n_steps)
{
    int m = 0;
    unsigned int seed = time(NULL) ^ omp_get_thread_num();
    srand(seed);

    for (int i = 0; i < n_steps; i++)
    {
        double x = (double) rand() / RAND_MAX;
        double y = (double) rand() / RAND_MAX;
        if (x * x + y * y < 1.)
        {
            m += 1;
        }
    }

    double pi = 4.0 * m / n_steps;
    printf("Estimated value of pi = %f\n", pi);
}
```

Problem 2

Here, I implemented a OpenMP version of Monte Carlo. I used `reduction` to avoid **data race** for `m`.

```
void compute_pi_monte_carlo_parallel(int n_steps)
{
    int m = 0;
    #pragma omp parallel
    {
        unsigned int seed = time(NULL) ^ omp_get_thread_num();

        #pragma omp for reduction(+:m)
        for (int i = 0; i < n_steps; i++)
        {
            double x = (double) rand_r(&seed) / RAND_MAX;
            double y = (double) rand_r(&seed) / RAND_MAX;
            if (x * x + y * y < 1.)
            {
                m++;
            }
        }
    }

    double pi = 4.0 * m / n_steps;
    printf("Estimated value of pi = %f\n", pi);
}
```

Test Wall Time Elapsed for Our Programs

To test how good our parallel program performs, we need to measure the **Wall Time** (Time consumed in the real world) of our program instead of CPU Time. For the measurement of this, the recommended way is the Linux/MacOS command `time` or the function `omp_get_wtime` provided by OpenMP other than `clock`.

I use the following code in my code for the measurement purpose (use `clock` if you wish to compile without OpenMP):

```

int n_steps = 100000000;
double start = omp_get_wtime();
...Your Code/Function For Parallel Computation
double end = omp_get_wtime();
double time_spent = end - start;
printf("Wall time taken: %f seconds\n", time_spent);

```

The testing result when we set `n_steps=100000000` :

```

# Parallel version
Estimated value of pi = 3.141505
Wall time taken: 0.229149 seconds
-----
# Nonparallel version
Estimated value of pi = 3.141704
Wall time taken: 1.495120 seconds

```

Problem 3

In this part we implement a MPI version of the Monte Carlo example.

First make sure we have installed OpenMPI properly:

```

mpicc --version # For compiling MPI program (We can also edit CMakeLists.txt to enable gcc to compile it)
mpiexec --version # For executing MPI program

```

Include `mpi.h` to make MPI functions visible to our code. In MPI program, `rank` is the id of the current thread that is doing the parallel tasks; `size` indicates how many threads are running in parallel. To boot up a MPI program, we use `MPI_Init(argc, argv)` at first.

Initialize MPI and get the parameters of our parallel program:

```

MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);

```

Similarly, we assign a random seed for each thread and use `rand_r()` for random number generation. We collect statistics respectively for each thread,

```

srand(rank);

for (int i = 0; i < amount; ++i) {
    const double x = (double) rand() / RAND_MAX;
    const double y = (double) rand() / RAND_MAX;

    if (x * x + y * y <= 1)
        ++m;
}

```

and reduce after all the jobs are done.

```

MPI_Reduce(&m, &total_hits, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);

if (rank == 0) {
    printf("PI is approximately equal to: %f\n", 4. * total_hits / n);
}

```

Put `MPI_finalize()` in the end of the MPI program to free the threads.

To execute MPI program with 4 threads, we use:

```

mpiexec -np 4 your_program_fullname

```

The result of our program:

```
Process 3 out of 4
Process 0 out of 4
Process 2 out of 4
Process 1 out of 4
Estimated Pi = 3.1414960000
```

Problem 4

See the pages after the Problem 6.

Problem 5

My implementation of parallel scalar product between vector and vector:

```
double a[5] = {1.0, 2.0, 3.0, 4.0, 5.0},
        b[5] = {1.0, 2.0, 3.0, 4.0, 5.0};
double res = .0;

#pragma omp parallel for reduction(+:res)
for (int i = 0; i < 5; i++) {
    res += a[i] * b[i];
}

printf("Dot product = %f\n", res);

return 0;
```

Suppose we use n processors for vector of order n , the time complexity of doing multiplication is $O(1)$ while the complexity of reduction is $O(\log n)$, when we *divide and conquer*. As a result,

$$T(n) = O(\log n).$$

Problem 6

My implementation of parallel dot product between matrix and vector:

```
// dot product mat and vec
double mat[2][2] = {{1.0, 2.0}, {3.0, 4.0}},
        vec[2] = {1.0, 2.0};
double res2[2] = {0.0, 0.0};
double start = omp_get_wtime();

// collapse(2) is used to collapse the two loops into one
#pragma omp parallel for collapse(2) reduction(+:res2[:2])
for (int i = 0; i < 2; i++) {
    for (int j = 0; j < 2; j++) {
        res2[i] += mat[i][j] * vec[j];
    }
}

double end = omp_get_wtime();
double time_spent = end - start;
printf("Wall time taken: %f seconds\n", time_spent);
printf("Matrix-vector product = [%f, %f]\n", res2[0], res2[1]);
```

Each row of the resulting vector is calculated in the same way with vector dot product, thus having a time complexity of $O(\log n)$. In the reduction phase, we still can reduce the results in parallel, which results in a $O(\log n)$ time complexity. What is different is that we need $n \cdot n$ processors for the reduction phase. Hence,

$$T(n) = O(\log n).$$

But....

This seems true. However, in EREW model, exclusive read is not allowed so we can not do n vector inner product at the same time so the time complexity should still be $O(n \log n)$.

Exercise 4

Preparations for the Proof (Lemma?)

Define the order of a binomial tree as the height of that binomial tree.

Observe the given binomial tree, we can discover that:

1. Binomial tree of order 0 has 1 nodes;
2. Binomial tree of order k is generated by putting two binomial tree of order $k - 1$ together by connecting their root;
3. Combining 1. and 2., we have that binomial tree of order k has exactly 2^k nodes;
4. Root of binomial tree of order k has exactly k children.




Figure 1: A binomial tree given by the Lecture Note

In the context of broadcast, we know that the total node of the binomial tree is p . It follows that the order of the binomial tree is $k = \log p$ (or we need some rounding up).




Figure 2: The smaller children is the same with the right children of the bigger children

If the smaller children are preferred, when the nodes start to send messages, as is shown on the Figure 1, we will process the children from left to right. We can note that the cost of processing the

smaller children is always smaller than the cost of bigger children (See Figure 2). So the subtask of broadcast for the smaller children will always finish earlier than the bigger children.

The statements above form the basis of our proof.

Proof

- First, it takes **$\log p$ steps** for the root node to inform all of its children of the message;
- Second, we find that the last children informed is of order $\log p - 1$, and it's manifest that it can not be processed in a instant (or $O(1)$). We need another $\log p - 1$ steps for this children to inform their children of the message;
- At last, it took $\log p + \log p - 1 + \dots + 1$ steps to make all the PEs informed, which sums up to $\log p(\log p + 1)/2$. This gives the time complexity $\Omega(\log^2 p)$. The Ω -notation arises because of the rounding up of $\log p$. ■