

# Exercises 1

---

Coauthored by: Teng, Chao & Jonas Karl, Böcker

## Problem 1

---

The nonparallel version of computing pi by Monte Carlo.

```
void compute_pi_monte_carlo(int n_steps)
{
    int m = 0;
    unsigned int seed = time(NULL) ^ omp_get_thread_num();
    srand(seed);

    for (int i = 0; i < n_steps; i++)
    {
        double x = (double) rand() / RAND_MAX;
        double y = (double) rand() / RAND_MAX;
        if (x * x + y * y < 1.)
        {
            m += 1;
        }
    }

    double pi = 4.0 * m / n_steps;
    printf("Estimated value of pi = %f\n", pi);
}
```

## Problem 2

---

Here, I implemented a OpenMP version of Monte Carlo. I used `reduction` to avoid **data race** for `m`.

```
void compute_pi_monte_carlo_parallel(int n_steps)
{
    int m = 0;
    #pragma omp parallel
    {
        unsigned int seed = time(NULL) ^ omp_get_thread_num();

        #pragma omp for reduction(+:m)
        for (int i = 0; i < n_steps; i++)
        {
            double x = (double) rand_r(&seed) / RAND_MAX;
            double y = (double) rand_r(&seed) / RAND_MAX;
            if (x * x + y * y < 1.)
            {
                m++;
            }
        }
    }

    double pi = 4.0 * m / n_steps;
    printf("Estimated value of pi = %f\n", pi);
}
```

## Test Wall Time Elapsed for Our Programs

---

To test how good our parallel program performs, we need to measure the **Wall Time** (Time consumed in the real world) of our program instead of CPU Time. For the measurement of this, the recommended way is the Linux/MacOS command `time` or the function `omp_get_wtime` provided by OpenMP other than `clock`.

I use the following code in my code for the measurement purpose (use `clock` if you wish to compile without OpenMP):

```

int n_steps = 100000000;
double start = omp_get_wtime();
...Your Code/Function For Parallel Computation
double end = omp_get_wtime();
double time_spent = end - start;
printf("Wall time taken: %f seconds\n", time_spent);

```

The testing result when we set `n_steps=100000000` :

```

# Parallel version
Estimated value of pi = 3.141505
Wall time taken: 0.229149 seconds
-----
# Nonparallel version
Estimated value of pi = 3.141704
Wall time taken: 1.495120 seconds

```

## Problem 3

---

In this part we implement a MPI version of the Monte Carlo example.

First make sure we have installed OpenMPI properly:

```

mpicc --version # For compiling MPI program (We can also edit CMakeLists.txt to enable gcc to compile it)
mpiexec --version # For executing MPI program

```

Include `mpi.h` to make MPI functions visible to our code. In MPI program, `rank` is the id of the current thread that is doing the parallel tasks; `size` indicates how many threads are running in parallel. To boot up a MPI program, we use `MPI_Init(argc, argv)` at first.

Initialize MPI and get the parameters of our parallel program:

```

MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &size);

```

Similarly, we assign a random seed for each thread and use `rand_r()` for random number generation. We collect statistics respectively for each thread,

```

srand(rank);

for (int i = 0; i < amount; ++i) {
    const double x = (double) rand() / RAND_MAX;
    const double y = (double) rand() / RAND_MAX;

    if (x * x + y * y <= 1)
        ++m;
}

```

and reduce after all the jobs are done.

```

MPI_Reduce(&m, &total_hits, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);

if (rank == 0) {
    printf("PI is approximately equal to: %f\n", 4. * total_hits / n);
}

```

Put `MPI_finalize()` in the end of the MPI program to free the threads.

To execute MPI program with 4 threads, we use:

```

mpiexec -np 4 your_program_fullname

```

The result of our program:

```
Process 3 out of 4
Process 0 out of 4
Process 2 out of 4
Process 1 out of 4
Estimated Pi = 3.1414960000
```

## Problem 4

---

See the pages after the Problem 6.

## Problem 5

---

My implementation of parallel scalar product between vector and vector:

```
double a[5] = {1.0, 2.0, 3.0, 4.0, 5.0},
        b[5] = {1.0, 2.0, 3.0, 4.0, 5.0};
double res = .0;

#pragma omp parallel for reduction(+:res)
for (int i = 0; i < 5; i++) {
    res += a[i] * b[i];
}

printf("Dot product = %f\n", res);

return 0;
```

Suppose we use  $n$  processors for vector of order  $n$ , the time complexity of doing multiplication is  $O(1)$  while the complexity of reduction is  $O(\log n)$ , when we *divide and conquer*. As a result,

$$T(n) = O(\log n).$$

## Problem 6

---

My implementation of parallel dot product between matrix and vector:

```
// dot product mat and vec
double mat[2][2] = {{1.0, 2.0}, {3.0, 4.0}},
        vec[2] = {1.0, 2.0};
double res2[2] = {0.0, 0.0};
double start = omp_get_wtime();

// collapse(2) is used to collapse the two loops into one
#pragma omp parallel for collapse(2) reduction(+:res2[:2])
for (int i = 0; i < 2; i++) {
    for (int j = 0; j < 2; j++) {
        res2[i] += mat[i][j] * vec[j];
    }
}

double end = omp_get_wtime();
double time_spent = end - start;
printf("Wall time taken: %f seconds\n", time_spent);
printf("Matrix-vector product = [%f, %f]\n", res2[0], res2[1]);
```

Each row of the resulting vector is calculated in the same way with vector dot product, thus having a time complexity of  $O(\log n)$ . In the reduction phase, we still can reduce the results in parallel, which results in a  $O(\log n)$  time complexity. What is different is that we need  $n \cdot n$  processors for the reduction phase. Hence,

$$T(n) = O(\log n).$$

But....

This seems true. However, in EREW model, exclusive read is not allowed so we can not do  $n$  vector inner product at the same time so the time complexity should still be  $O(n \log n)$ .