

## Exercise 1 (Unimplemented parts)

We didn't implement the Parallel Partition successfully because of some technical details. But I will give the pseudocode of the parallel partition below. It may look verbose and complicated, but the key point is to find the correct index to insert the elements. (See line 32. - 34.)

### Parallel Partition

```
1: procedure PARALLELPARTITION( $A, n, \text{pivot}, \text{offset}$ )
2:   ▷ The following 3 arrays are used for storing the count of element in each processors, that are
   smaller/equal/larger to pivot.
3:   localSmaller  $\leftarrow$  Array[1..p]
4:   localEqual  $\leftarrow$  Array[1..p]
5:   localBigger  $\leftarrow$  Array[1..p]
6:   ▷  $B$  as a temporary array to store the rearranged array
7:    $B \leftarrow$  Array[1..n]
8:   for tid=1..p in parallel do
9:     smallerCount  $\leftarrow$  0
10:    equal  $\leftarrow$  0
11:    biggerCount  $\leftarrow$  0
12:    for j=start(tid)...end(tid) do
13:      if  $A[j + \text{offset}] < \text{pivot}$  then
14:        smallerCount ++
15:      else if  $A[j + \text{offset}] == \text{pivot}$  then
16:        equal ++
17:      end
18:      else
19:        biggerCount ++
20:      end
21:    end
22:    end
23:    localSmaller[tid]  $\leftarrow$  smallerCount
24:    localEqual[tid]  $\leftarrow$  equal
25:    localBigger[tid]  $\leftarrow$  biggerCount
26:  end
27:  prefixSmaller  $\leftarrow$  pPrefAdd(localSmaller)
28:  prefixEqual  $\leftarrow$  pPrefAdd(localEqual)
29:  prefixBigger  $\leftarrow$  pPrefAdd(localBigger)
30:  ▷ Let prefix*[0] be 0
31:  for tid=1..p in parallel do
32:    sIdx  $\leftarrow$  prefixSmaller[tid - 1]
33:    eIdx  $\leftarrow$  prefixSmaller[p] + prefixEqual[tid - 1]
34:    bIdx  $\leftarrow$  prefixSmaller[p] + prefixEqual[p] + prefixBigger[tid - 1]
35:    for j=start(tid)...end(tid) do
36:      if  $A[j + \text{offset}] < \text{pivot}$  then
37:         $B[\text{sIdx} + +] \leftarrow A[j + \text{offset}]$ 
38:      else if  $A[j + \text{offset}] == \text{pivot}$  then
39:         $B[\text{eIdx} + +] \leftarrow A[j + \text{offset}]$ 
40:      end
41:      else
42:         $B[\text{bIdx} + +] \leftarrow A[j + \text{offset}]$ 
43:      end
44:    end
45:  end
46:  end
47:  Copy  $B$  back to  $A[\text{offset}..\text{offset} + n]$ 
48: end
```

## Exercise 3

### 3.1

The **insert** operation can result in concurrent write if it is not carefully designed, given the case where  $\text{index}(A[i]) == \text{index}(A[j])$ .

The other parts are guaranteed to be thread-safe.

### 3.2

For a parallel version of **copy**( $A, B$ ), we will need the help of another array  $C$ .

```
1: procedure PARALLEL_COPY( $A, B$ )
2:   ▷  $C$  stores the size of each bucket in  $B$ 
3:    $C \leftarrow \text{Array}[1..nb]$ 
4:   pfill( $C, 0$ )
5:   for  $i=1..nb$  in parallel do
6:      $C[i] \leftarrow \text{len}(B[i])$ 
7:   end
8:    $P \leftarrow \text{pPrefAdd}(C)$ 
9:   for  $i=1..nb-1$  in parallel do
10:    ▷ we set  $P[0]$  as 0
11:     $\text{beg} \leftarrow P[i-1] + 1$ 
12:     $\text{end} \leftarrow P[i]$ 
13:    for  $j=\text{beg}..\text{end}$  in parallel do
14:       $A[j] \leftarrow B[i][j - \text{beg}]$ 
15:    end
16:  end
17: end
```

#### 3.2.1 Comments

- For the beg and end in the above algorithm, we note that  $(P[i] - (P[i-1] + 1) + 1)$  is equal to  $C[i]$ , namely  $\text{len}(B[i])$ .
- I don't see the necessity of using **pfill**( $\cdot$ ). It's likely used for avoiding data corruption.

#### 3.2.2 Complexity Analysis

<i>Operation</i>	<i>Time Complexity</i>	<i>Processors Required</i>
Get the lengths	$O(1)$	nb
Parallel Prefix Sum	$O(\log nb)$	nb/2
Insert	$O(1)$	n
<b>Overall</b>	$O(\log nb)$	n

For fast Parallel Bucket Sort, we need enough number of buckets. However, when it comes to the parallel copy operation, a lower nb is preferred. Herein lies a tradeoff.

## Exercise 5

### (a)

Given an index  $i$ , to find the final position  $j$  of  $a_i$  in the merged list, we can use Binary Search to find the biggest element  $b_k$  in  $B$  that is smaller than  $a_i$  (or the smallest element in  $B$  that is bigger

than  $a_i$ ), which requires  $O(\log n)$  comparisons. (See section **Parallel Merge sort** in the lecture slide for example) Then,  $j = i + k$ , as there are  $i + k$  elements in total that are smaller than  $a_i$ .

### (b)

Denote by **computeInsertPos**( $A, B, i$ ) the algorithm described in (a), which returns the final position  $j$  to insert our  $a_i = A[i]$ .

```

1: procedure PARALLEL_MERGE( $A, B$ )
2:    $C \leftarrow \text{Array}[1..2n]$ 
3:   for  $i = 1..n$  in parallel do
4:      $\triangleright j_1 \neq j_2$  as the elements are unique
5:      $j_1 \leftarrow \text{computeInsertPos}(A, B, i)$ 
6:      $j_2 \leftarrow \text{computeInsertPos}(B, A, i)$ 
7:      $C[j_1] \leftarrow A[i]$ 
8:      $C[j_2] \leftarrow B[i]$ 
9:   end
10: end

```

#### (b.1) Comments

- If the elements are not unique and  $j_1 == j_2$  by accident, we can set  $j_2$  as  $j_1 + 1$  to avoid data race.
- The total work  $T_1$  is apparently  $O(n \log n)$  and the span  $T_\infty$  is  $O(\log n)$ .

### (c)

This task seems impossible from the software perspective. With  $O(n)$  processors, we can compare  $a_i$  with  $b_{1..n}$  in  $O(1)$ . (but don't we need to broadcast  $a_i$  to all the processors?)

That said, it's still not possible to reduce the results in  $O(1)$ . Consider the following 2 cases:

1.  $\max_{\{k: a_i \geq b_k\}} k;$
2.  $\sum_{\kappa \in \{k: a_i \geq b_k\}} 1.$

In both cases, the reduction takes  $O(\log \#\{k : a_i \geq b_k\})$ .  $O(1)$  is only made possible if we have some hardware "hacks" or some magic.

If such a hack exists, we can use  $n \cdot O(n) = O(n^2)$  processors for merging two lists in  $O(1)$ .